

---

# SUMMARY OF DATABASE REPLICATION: A TUTORIAL

By S. Xiao Fernández Marín

## 1 Summary

Why it is needed several copies of the data?

- **scalability:** things go faster, depending on how many clients  
r(requests,, easier using hash)  
w(update data,, harder bc conflicts)
- **fault tolerance:** high availability, if one server crashes, you can redirect to another server if some fails
- **decreasing response time:**  
**geographic distribution),** the latency is limited by speed of light so there is a long response time if the server is far (norway to NZ)  
**off-line (disconnected) operation** e.g: google docs: work offline and then when connecting to internet, update.

Terms:

- **Read-one-write-all-available (ROWAA)** NOT ALWAYS IMMEDIATELY
  - Reads executed on one node - enough only read from one node to get the value
  - Writes executed on all nodes - all the nodes have updates
- **Transaction location**
  - Primary - all the updates in that server or part of the db or all the db,, distribute the weight of primary
  - Update any were - clients update on any server
- **Synchronization strategy**
  - Eager: early on propagating updates to other servers,, to commit a transaction we have to replicate all the updates to all the servers first. *Coordinate early* and *Can guarantee strong consistency*
  - Lazy: not propagate before, we can do it later. *Don't coordinate* and *Weak consistency*: Stale, Temporarily inconsistent.
- **1-copy semantics:** single centralize system, we assume that we have one single server even if we have more that one server
  - Behave as if there was one centralized copy of the data: No distribution,, clients doesn't see the replication, they see an equivalent of a centralized system... we want the system to behave like this
  - Atomic operations: No in-between states;; no full transactions, just atomic
- **Two-phase commit (2PC):** 1 phase commit: one decide other do way of agree as a DS to abort or commit an atomic transaction and allow nodes to detect deadlocks without having to coordinate with other servers
  - Distributed agreement on whether to commit or abort
  - coordinator: not the same one for all the transactions but the same one for the same transaction.. if there are more, behave as 1,, gather votes from all the participants
  - participators- $i$  can decide to commit or abort- if someone says to abort, it aborts. to commit, everyone agrees to commit

Depending on the transaction location and the synchronization strategy

---

		transaction location: WHERE?	
		primary copy	update anywhere
synchronization point: WHEN?	eager	+ simple cc + strong consistency + potentially long response times - inflexible	+ flexible - complex cc
	lazy	+ simple cc + often fast - stale data - inflexible	+ flexible + always fast - inconsistency - conflict resolution

Figure 1: Categories

- **Eager primary copy** update every w, can update in central node
  - All read/write transactions on primary node
    - Local reads,, fast
    - Writes are distributed to all nodes,, can be bottleneck, we have to send them to the primary node
  - Read-only transactions on any node
  - Nearly the same as centralized system
    - Globally serializable
    - Applications must know if transactions are r/w or read-only
  - Updates are expensive
- **Eager update anywhere** update every w, can update in all nodes
  - All transactions can happen anywhere, no need to differentiate between the transactions of r-only or r/w
    - Same semantics as centralized system
      - Globally serializable
      - Applications don't have to know it's distributed
    - Distributed deadlocks may happen bc update in any node and we can read from another
    - Both reads and updates are expensive but very flexible
- **Lazy primary copy** update later, can update in central node
  - All read/write transactions on primary node
    - Local reads
    - Local writes (sent to other nodes post commit) – someone reads old info bc central node didn't update the other ones
  - Read-only transactions on any node
  - Low overhead
  - Weaker consistency
    - Stale reads
  - Weaker durability in case of primary copy fail during commit
- **Lazy update anywhere** update later, can update in all nodes
  - All transactions can happen anywhere
    - Low overhead
    - Durability not guaranteed
    - Atomicity may be lost

- 
- Weak consistency and isolation
    - Needs conflict resolution in case of concurrent updates; one node can do an update and another one other in the same bank account eg

### **Eventual consistency**

- Temporary inconsistencies are allowed
  - May even be the normal state
- In a steady state, consistency is achieved
  - Left alone
  - Without any updates
  - Eventually
- Consistent values vs. consistent history
  - Different histories can lead to the same values

### **Session consistency**

- A client first issues read/write transaction T1, then read-only transaction T2
- Will T2 see T1's changes?
  - Lazy replication may not have propagated fully
- Global transaction IDs
  - Freshness requirements on read replicas
- Requires application/driver support

Other parameters

- **Cluster vs. WAN**
  - Cluster, everything close together
    - Low latency
    - High bandwidth
  - WAN, more far away
    - High latency
    - Low bandwidth
    - Geographic partitioning?
- **Statement vs. object replication**
  - Statement replication, mandas la funcion
    - E.g., SQL statement
    - Executed separately on each node
    - Efficient when statement is smaller than changed data
    - Semantically difficult, e.g., RAND()
  - Object replication, mandas el objeto
    - Executed once
    - Changed data is replicated
    - Efficient when changes are small
- **Concurrency control**
  - Pessimistic: **Locking**
    - leads to possible distributed deadlocks conflicts
  - Snapshot isolation
    - Only write/write conflicts
  - Optimistic: Timestamps
    - No deadlocks
    - Centralized or distributed validation

---

- **Architecture**

- Kernel-based, replication in the core
  - Built-in replication
  - White box
- Middleware
  - Black box, helps replication
  - Gray box, some db primitives replication with help of special APIs
  - Centralized or distributed middleware layer

- **Replication hierarchies**

- Write vs. read
  - Replication doesn't give scalability if 100- 100% reads is infinitely scalable
- Fan-out
  - How many secondaries to each primary?
  - Replication is not free
- How many levels?
  - Lazy replication lag increases with each level