
SUMMARY OF NoSQL DATABASES AND BIG DATA STORAGE SYSTEMS

By S. Xiao Fernández Marín

1 Summary

The idea of NoSQL is to store large amount of data and be efficient with working with that data

- Availability more important than consistency: Inspired by the CAP theorem
- It's hard to guarantee ACID: Move implementation cost from server (DB system) to application (client)
- It's complex to implement SQL: Move complex operations from server to application and Imperative languages avoids implementing optimizer
- Schemaful forces application developers to plan ahead: Move schema enforcement from server to application, logic in the app side

Typical characteristics

- Horizontal scalability
- High availability
- Weak consistency
 - Eventual consistency
 - Applications must handle inconsistencies
- Data distribution
 - Horizontal sharding (each row of a table in different server) bc horizontal scalability, the app handle the app, so the app combines the sharding // dist db partitioning: do query independently where the data is located and db will locate
 - Limited cross-shard functionality
- Simple data model
 - Key/value, like hash
 - Limited indexing
- Schemaless, no specific requirements in schema
 - Dynamic typing
 - Schema encoded in application logic
- Data model tied closer to programming language
 - Object model, JSON from Java script eg
- Denormalized data models,, //NOT SURE duplicated data
- Less advanced query languages
 - Create, read, update, delete (CRUD)
 - Syntax often close to source code
 - Imperative instead of declarative
 - Join often on application side
- No transactions
 - No ACID properties
 - Little isolation

Simpler system, data models:

- Key-value databases: Value is a binary string or slightly more advanced
- Document databases: Values are objects (often compatible with JSON)
- Column databases: Similar to RDBMSs, but weaker schema
- Graph databases: Nodes and links with attributes (similar to ER models)

Key-value databases Simples of them all

A key points to a value, hash

- Unique key
- Value is string of bytes
- Values is an object (often JSON compatible)

Very simple data model

No complex operations

- CRUD
- CRD (no updates)

Examples: Amazon DynamoDB, Voldemort, Oracle NoSQL Database, Redis

Document databases

Really object databases

- Structured objects like JSON, not documents like HTML

Similar to key-value stores

- Values are objects, not random byte strings
- Richer query capabilities

Usually (almost) JSON compatible

Key can be separate or embedded in value object

Indexing on non-key members, not like key-value

Examples: MongoDB, CouchDB, RethinkDB, Couchbase

Column databases

More similar to RDBMSs

- Tabular data model
- More SQL like query languages

Weak schema

- Columns not predefined, may vary between rows
- Predifined "column families"

Often versioned cells

Sometimes called "wide column stores"

Examples: Google Bigtable, Apache Hbase

Graph databases

Stores vertexes and edges

- Often called nodes and relationships
- Both vertexes and edges may have attributes

Close to (E)ER model

Searched/navigated with path expressions

Special query languages

Examples: Neo4j, RedisGraph, Amazon Neptune

NoSQL in RDBMSs

- All major RDBMSs have some NoSQL functionality
 - JSON object stores
 - Graph functionality
- RDBMS replication
 - High availability
 - Sharding
- Implement superset of NoSQL functionality
 - Optimized for different use cases

NoSQL vs. SQL

- **NoSQL** systems are becoming more advanced: query languages, Joins and Transactions
- **SQL** systems are adding more schemaless features: - JSON: Similar to object-relational and XML (in SQL since year 2000)
 - SQL/JSON standard
 - Graph query languages